



Brunel
University
London

CS5810 High Performance Computational
Infrastructures
Coursework for 2022/23

Student ID: 2242297
Academic Year: 2022-23

1. Introduction, Problem Description & Associated Dataset

1.1 Introduction

Comcast Corporation Pvt Ltd (imaginary) is a Telecommunication organization that provides telecommunication services to millions of customers. Typically, a single customer can generate a significant amount of data, including internet usage, TV viewing habits, and customer churn rate which requires advanced data processing and high-performance computational infrastructure capabilities to manage and make sense of the collected data. To analyze sales transactions, Comcast Telecommunication uses an advanced data analytics system that utilizes Hadoop and machine learning algorithms. The system can identify patterns in customer behavior, such as it provides insights into which customers are likely to churn, enabling the company to take proactive measures to retain them. Due to the vast amount of data being generated every day, Comcast uses the Hadoop system, and cloud storage to store the data. The company employs a robust security system to ensure the privacy and security of the data.

1.2 Problem Description

The biggest problem of the telecommunication organization today is that customers discontinue using the company's services. When customers cancel their subscriptions, the company loses revenue, and it becomes more expensive to acquire new customers. High customer churn rates can lead to reduced revenue, increased customer acquisition cost, and damage to a company's reputation. To understand the factors that contribute to customer churn, the telecommunication industry has decided to do research on the below question.

Research Question:

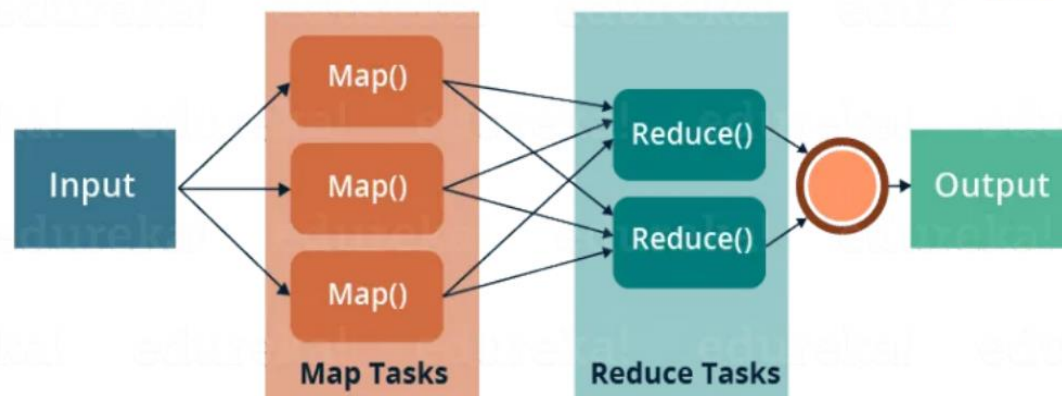
1. "Are the customers canceling/leaving the subscription when they subscribe smaller contract period (Month-to-Month) and how does it vary from people having dependents?"

The question being asked is whether customers are more likely to cancel or leave their subscriptions when they choose smaller contract periods, such as month-to-month, and whether this varies for customers with dependents. This research question is crucial as it can provide valuable insights into customer behavior and enable the industry to take proactive measures to reduce customer churn and improve customer retention.

Customer satisfaction is a key component of any industry, and understanding whether customers are more likely to cancel their subscriptions when they opt for smaller contract periods can help the telecommunication industry develop targeted retention strategies. For example, they could offer incentives to customers who sign up for longer contract periods or create customized plans that cater to the needs of customers with dependents. By implementing such strategies, the industry can improve customer satisfaction and decrease churn rates.

This research question requires the expertise of a senior data scientist who can analyze customer data to identify patterns and trends. The data engineering team also plays a crucial role in the process as they are responsible for managing the data infrastructure and ensuring that the data is accurate and reliable. Additionally, the marketing team may be involved in creating retention strategies based on the findings of the research.

What is MapReduce?



What is MapReduce - MapReduce Tutorial

MapReduce is a programming framework that enables us to carry out distributed parallel processing on huge datasets in a distributed environment. Map and reduce are two separate activities that makeup 'MapReduce'. It is a two-part process, consisting of mapping and reducing activities. The mapper phase is the initial task, where a block of data is read and processed to generate key-value pairs in an intermediate output. After the mapper phase, the Reducer phase follows, which is indicated by the name "MapReduce."

The Reducer receives key-value pairs from multiple map jobs, and its purpose is to condense the larger set of intermediate data tuples or key-value pairs into a smaller set of intermediate data tuples. The input to the reducer is composed of key-value pairs generated from the output of the Mapper, and the output of the Reducer is also in the form of key-value pairs.

Now, Why MapReduce in this research?

1. The varied range of data created by the telecommunications industry is well-suited for MapReduce since it can handle a wide range of data kinds and formats. Also, it has built-in fault tolerance, which means that it will continue to work even if some of the nodes in the cluster are down/fail.
2. MapReduce operations will effectively meet the business demand in this situation since the chunk of data is too large to be handled by any other traditional data warehousing or RDBMS systems.
3. Since this is a batch-processing job situation that calls for contract-based data analysis and report generation, building a MapReduce job is the best way to meet this need.

1.3 Associated Dataset

The dataset for this analysis was obtained from Kaggle.com and can be found at the following URL, <https://www.kaggle.com/blatchar/telco-customer-churn>. The dataset contains a total of 21 columns and 7,043 rows. The size of the dataset is around 955KB. It has information about customer demographics such as age, gender, dependency, and usage patterns, and account information such as the type of plan and the contract period along with whether each customer churned or remained with the company. The complete data dictionary is mentioned below in the table.

Data Dictionary:

Column Name	Description	Datatype
customerID	Customer identification Number	String
gender	Whether the customer is a male or a female	String
SeniorCitizen	Whether the customer is a senior citizen or not	Integer
Partner	Whether the customer has a partner or not	Boolean
Dependents	Whether the customer has dependents or not	Boolean
tenure	Number of months the customer has stayed with the company	Integer
PhoneService	Whether the customer has a phone service or not	Boolean
MultipleLines	Whether the customer has multiple lines or not	String
InternetService	Customer's internet service provider	String
OnlineSecurity	Whether the customer has online security or not	String
OnlineBackup	Whether the customer has online backup or not	String
DeviceProtection	Whether the customer has device protection or not	String
TechSupport	Whether the customer has tech support or not	String
StreamingTV	Whether the customer has streaming TV or not	String
StreamingMovies	Whether the customer has streaming movies or not	String
Contract	The contract term of the customer	String
PaperlessBilling	Whether the customer has paperless billing or not	Boolean
PaymentMethod	The customer's payment method	String
MonthlyCharges	The amount charged to the customer monthly	Float
TotalCharges	The total amount charged to the customer	Float
Churn	Whether the customer churned or not	Boolean

2. Design and Implementation

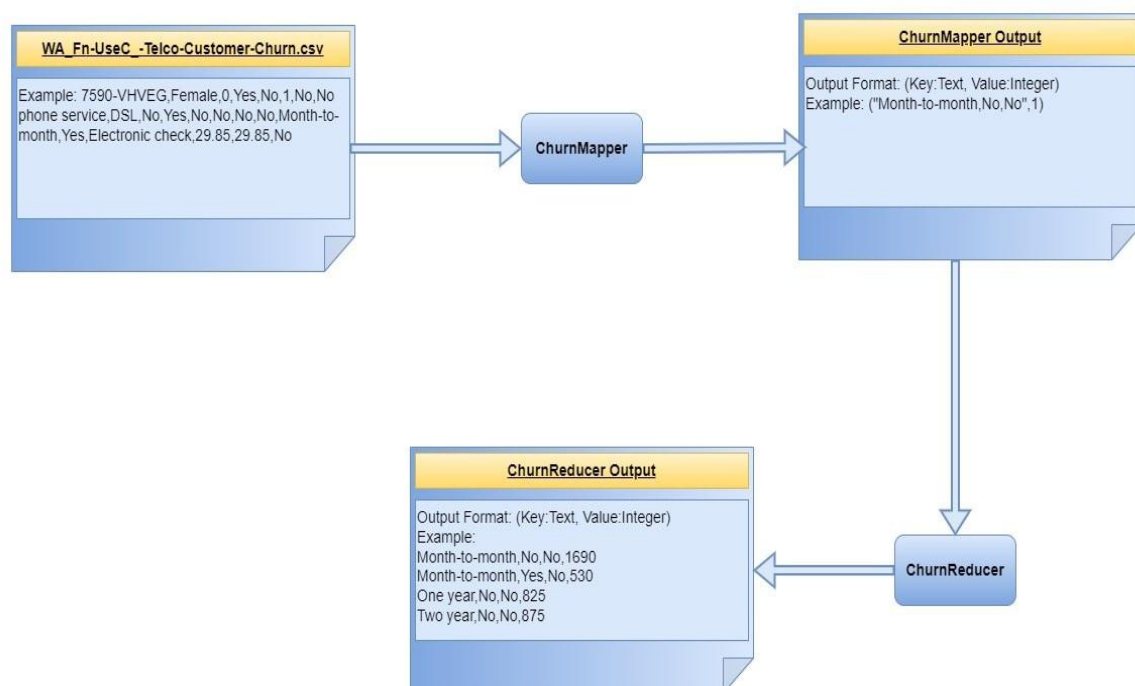
2.1 Design:

To answer the question, the data in the customer churn CSV file must be filtered to get contract, dependents, and churn information of all customers and then the logic to find out the number of customers churned based on contract and how it varies from people having dependents. So, I used a single MapReduce Job to achieve this. First data will be passed through the mapper class (ChurnMapper) to get the fields required to filter out

from customer information. After that, a reducer class (ChurnReducer) takes the output of the mapper class and then the business logic is applied on it. A driver class (ChurnAnalysis) is implemented to run the whole project(.jar).

Component	Input	Input Types	Output	Output Types
Churn Mapper	Key: Line Number Value: Complete Line Text from WA_Fn-UseC_-Telco-Customer-Churn.csv file.	K: LongWritable V: Text	Key: (Contract, Dependents, Churn) Value: 1 Ex: <i>((Month-to-month,No,No), 1)</i>	K: Text V: IntWritable
Churn Reducer	Key: (Contract, Dependents, Churn) Value: (1, 1, 1,...)	K: Text V: List of Intwritables	Contract, Dependents, Churn, no. of customers churn Ex: <i>Month-to-month,No,No,1690</i>	K: Text V: Intwritable

Data Flow:



2.2 Implementation Steps:

1. As per the MapReduce programming model, create the below three classes and generate a .jar file for this requirement.

- a) ChurnMapper
- b) ChurnReducer
- c) ChurnAnalysis

- The complete implementation and the code are in the separate file '**Mapreduce_ChurnAnalysis**' and the file has been submitted along with this coursework.

2. Firstly, start the HDFS file system using commands 'start-dfs.sh' and 'start-yarn.sh' and confirm if all the services are running or not using a command called 'jps'

```
(base) hadoop@hadoop-VirtualBox:~$ start-dfs.sh
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [hadoop-VirtualBox]
(base) hadoop@hadoop-VirtualBox:~$ start-yarn.sh
Starting resourcemanager
Starting nodemanagers
(base) hadoop@hadoop-VirtualBox:~$ jps
4113 Jps
3187 DataNode
3797 NodeManager
3015 NameNode
3389 SecondaryNameNode
3566 ResourceManager
(base) hadoop@hadoop-VirtualBox:~$
```

Figure 1

- As per figure 1, all the services are up and running.

3. Create the directory in the HDFS file system to provide our input to the HDFS and execute the MapReduce job. In this case, I have created the 'Telecom' directory and then created the 'Input' directory inside the 'Telecom' directory as shown below.

```
hadoop@hadoop-VirtualBox: ~
(base) hadoop@hadoop-VirtualBox:~$ hdfs dfs -mkdir /Telecom
(base) hadoop@hadoop-VirtualBox:~$ hdfs dfs -mkdir /Telecom/Input
(base) hadoop@hadoop-VirtualBox:~$
```

Figure 2

4. Copy and place the input file 'WA_Fn-UseC_-Telco- Customer-Churn.csv' in the 'Input' folder using the 'put' command as shown below.

Note: Here, I have renamed the original downloaded file from 'WA_Fn-UseC_-Telco-Customer-Churn' to 'Telco_Customer_Churn'.

```
hadoop@hadoop-VirtualBox: ~  
(base) hadoop@hadoop-VirtualBox:~$ hdfs dfs -mkdir /Telecom  
(base) hadoop@hadoop-VirtualBox:~$ hdfs dfs -mkdir /Telecom/Input  
(base) hadoop@hadoop-VirtualBox:~$ hdfs dfs -put Telco_Customer_Churn.csv /Telecom/Input  
(base) hadoop@hadoop-VirtualBox:~$
```

Figure 3

5. Now, check whether the input file 'Telco_Customer_Churn.csv' is placed inside the /Telecom/Input directory or not. To do that, login to the interface, by default, it is available at <http://localhost:9870>. Next, go to Utilities and browse the file system '/Telecom/Input/'. Notice that the file 'Telco_Customer_Churn' is inserted in /Telecom/Input/ directory.

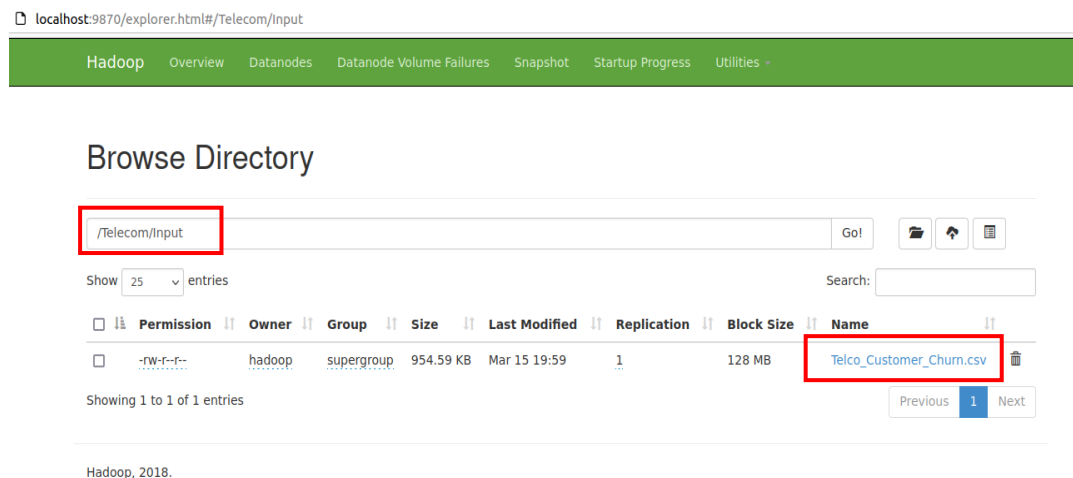


Figure 4

6. Now, In the Eclipse IDE, create the project and the required Java classes and export the .jar file. (Steps to create the java classes and exporting the .jar file mentioned in section 2.3 JAVA Classes Implementation and 2.4 Export the Jar file respectively)
7. Finally, run the exported jar file using the below command.

hadoop jar <jar file name> <driver class name> <input path> <output path>

Ex:<hadoop jar Downloads/CustomerChurn.jar org.myorg.ChurnAnalysis
/Telecom/Input /Telecom/Output>

```
hadoop@hadoop-VirtualBox: ~  
(base) hadoop@hadoop-VirtualBox:~$ hadoop jar Downloads/CustomerChurn.jar org.myorg.ChurnAnalysis /input /Telecom/out  
put  
2023-03-15 19:31:33,940 INFO impl.MetricsConfig: loaded properties from hadoop-metrics2.properties  
2023-03-15 19:31:34,024 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).  
2023-03-15 19:31:34,024 INFO impl.MetricsSystemImpl: JobTracker metrics system started  
2023-03-15 19:31:34,087 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Impleme  
nt the Tool interface and execute your application with ToolRunner to remedy this.  
2023-03-15 19:31:34,163 INFO input.FileInputFormat: Total input files to process : 1  
2023-03-15 19:31:34,182 INFO mapreduce.JobSubmitter: number of splits:1  
2023-03-15 19:31:34,301 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local326450544_0001  
2023-03-15 19:31:34,302 INFO mapreduce.JobSubmitter: Executing with tokens: []  
2023-03-15 19:31:34,409 INFO mapreduce.Job: The url to track the job: http://localhost:8080/  
2023-03-15 19:31:34,410 INFO mapreduce.Job: Running job: job_local326450544_0001  
2023-03-15 19:31:34,415 INFO mapred.LocalJobRunner: OutputCommitter set in config null  
2023-03-15 19:31:34,419 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2  
2023-03-15 19:31:34,419 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under ou  
tput directory:false, ignore cleanup failures: false  
2023-03-15 19:31:34,420 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapreduce.lib.output.FileOut  
putCommitter  
2023-03-15 19:31:34,464 INFO mapred.LocalJobRunner: Waiting for map tasks  
2023-03-15 19:31:34,464 INFO mapred.LocalJobRunner: Starting task: attempt_local326450544_0001_m_000000_0  
2023-03-15 19:31:34,495 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2  
2023-03-15 19:31:34,495 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under ou  
tput directory:false, ignore cleanup failures: false
```

Figure 5

8. When the job runs successfully, no error messages will be displayed after the completion of MapReduce as shown below.

```
hadoop@hadoop-VirtualBox: ~  
IO_ERROR=0  
WRONG_LENGTH=0  
WRONG_MAP=0  
WRONG_REDUCE=0  
File Output Format Counters  
Bytes Written=263  
2023-03-15 20:50:32,418 INFO mapred.LocalJobRunner: Finishing task: attempt_local1403515403_0001_r_000000_0  
2023-03-15 20:50:32,420 INFO mapred.LocalJobRunner: reduce task executor complete.  
2023-03-15 20:50:32,724 INFO mapreduce.Job: Job job_local1403515403_0001 running in uber mode : false  
2023-03-15 20:50:32,726 INFO mapreduce.Job: map 100% reduce 100%  
2023-03-15 20:50:32,728 INFO mapreduce.Job: Job job_local1403515403_0001 completed successfully  
2023-03-15 20:50:32,735 INFO mapreduce.Job: Counters: 35  
File System Counters  
FILE: Number of bytes read=10448  
FILE: Number of bytes written=943768  
FILE: Number of read operations=0  
FILE: Number of large read operations=0  
FILE: Number of write operations=0  
HDFS: Number of bytes read=1955002  
HDFS: Number of bytes written=263  
HDFS: Number of read operations=17  
HDFS: Number of large read operations=0  
HDFS: Number of write operations=4  
Map-Reduce Framework
```

Figure 6


```

hadoop@hadoop-VirtualBox: ~
Combine input records=7043
Combine output records=12
Reduce input groups=12
Reduce shuffle bytes=294
Reduce input records=12
Reduce output records=12
Spilled Records=24
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=0
Total committed heap usage (bytes)=549978112
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
Bytes Read=977501
File Output Format Counters
Bytes Written=263
(base) hadoop@hadoop-VirtualBox:~$

```

Figure 7

9. Check the output of the job which is created in the following location.
/Telecom/Output/part-r-00000

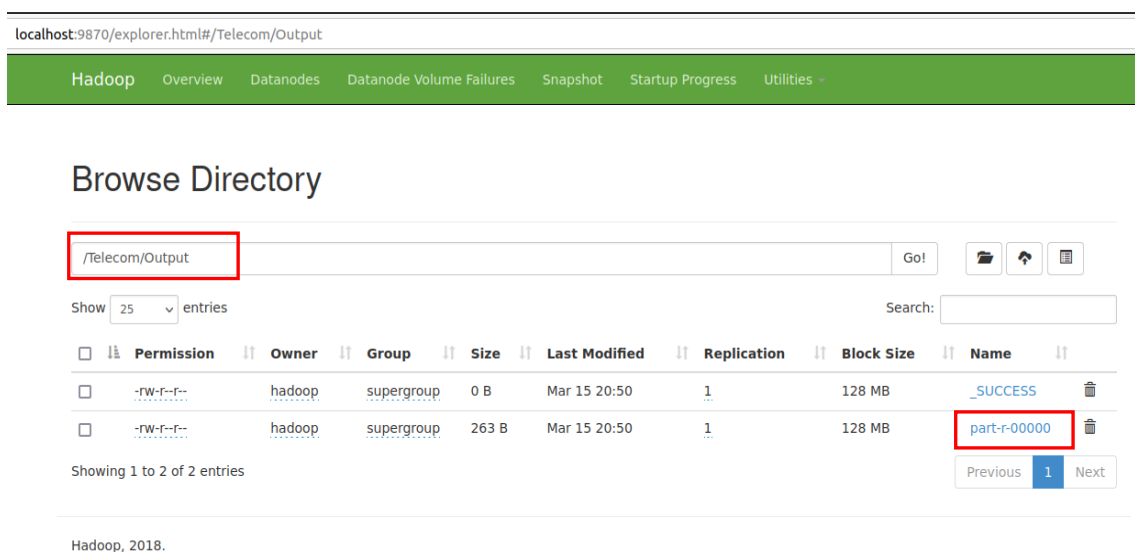


Figure 8

2.3 JAVA Classes Implementation

1. Create the project as shown in figure 9. Also, four additional jars were used in the development of this project.

- hadoop-common.jar – available at
</usr/local/hadoop/share/hadoop/common>
- hadoop-mapred-client-core.jar – available at
</usr/local/hadoop/share/hadoop/mapreduce>
- hadoop-mapred-client-jobclient.jar – available at
</usr/local/hadoop/share/hadoop/mapreduce>

- d. log4j.jar – available at
</home/hadoop/HPCI_Jars >

2. To obtain the log4j jar file needed for the project, you can download it from the official Apache log4j website at <https://logging.apache.org/log4j/1.2/download.html>. A separate file named "Mapreduce_ChurnAnalysis" has been provided with clear instructions for downloading the jar file. Log4j is a Java logging framework developed by the Apache Software Foundation and is used to track and manage program logs.

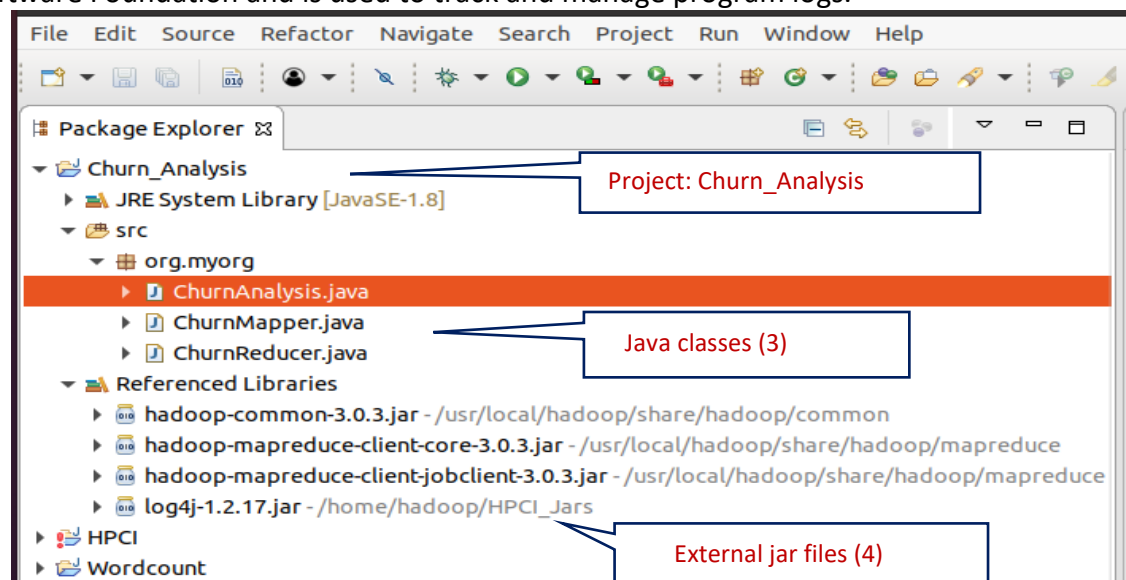


Figure 9

ChurnMapper (Mapper Class):

The Mapper class will take a line from the CSV file one by one as input. For each line, first, it will convert the line into a string and store it in a line variable. If the line is the header of data, then, it will avoid that line else it will split the line by comma to separate every field of a row. We need only contracts, dependents, and churn fields for this research question. So, these fields are accessed by using their respective indexes. The result of the mapper class consists of a key and a value. Mapper class will return contract, dependents and churn as the key in the form of Text and 1 as the value in the form of IntWritable. So, for each customer, we will get its contract, dependents, and churn as keys and 1 as values by using the mapper class.

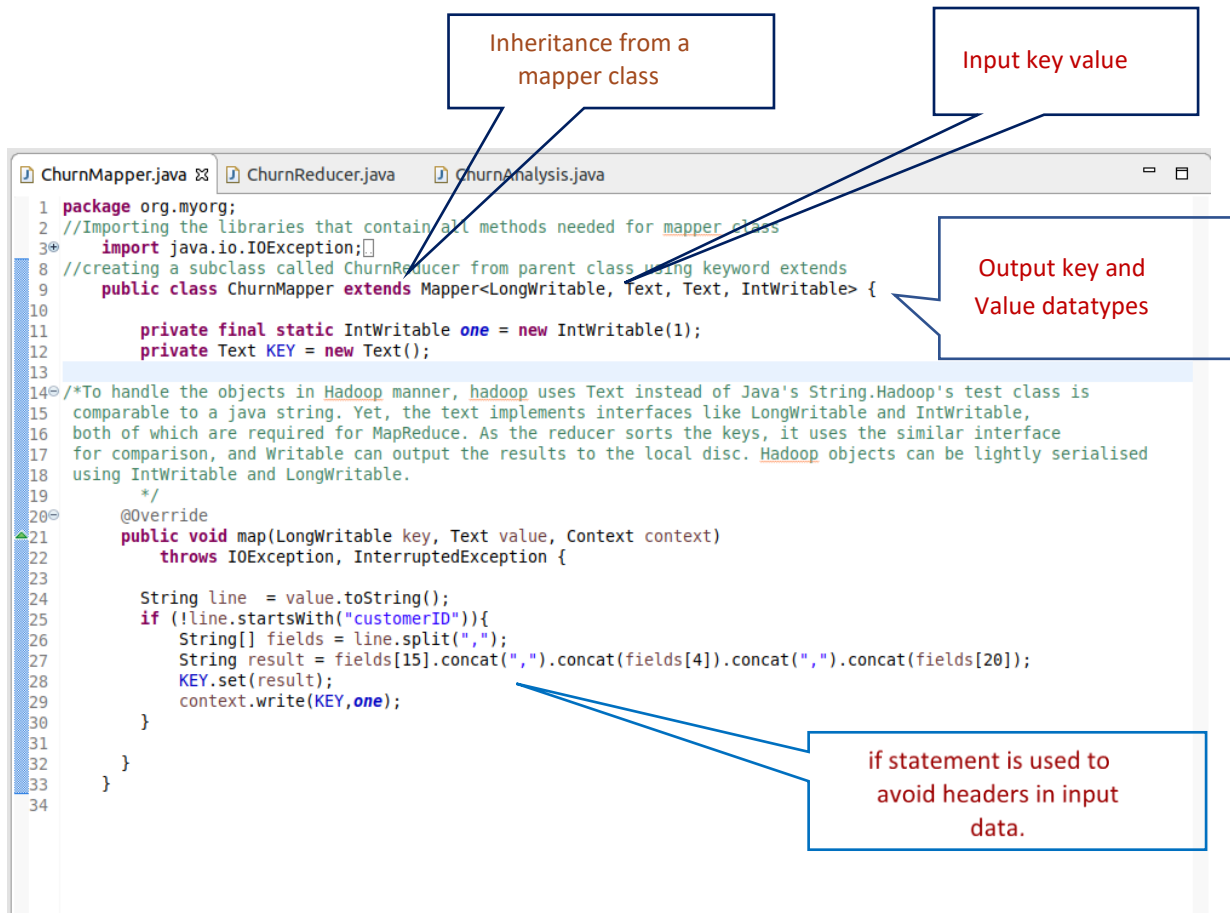


Figure 10

ChurnReducer (Reducer Class):

The Reducer class will take the results of the mapper class as input. It will group the values of the mapper class result by key and will aggregate those values. In this reducer class, the sum is applied as an aggregation function. For each key, it will take the 'sum' of all of its values. Reducer class will return contract, dependents, and churn as key in form of Text and 'sum of all values of that key' as the value in the form of IntWritable. So, the output would be in the following form:

Contract, dependent, churn, Sumofvalues

Example: Month-to-Month, No, Yes, 1396

In this example, 'Month-to-month' is a contract, 'No' is dependents, and 'Yes' is churn. Thus, 1396 is the total number of customers who have a Month-to-month contract, have 'No' dependents, and have churned the subscription (which means the customer has unsubscribed/left the contract).

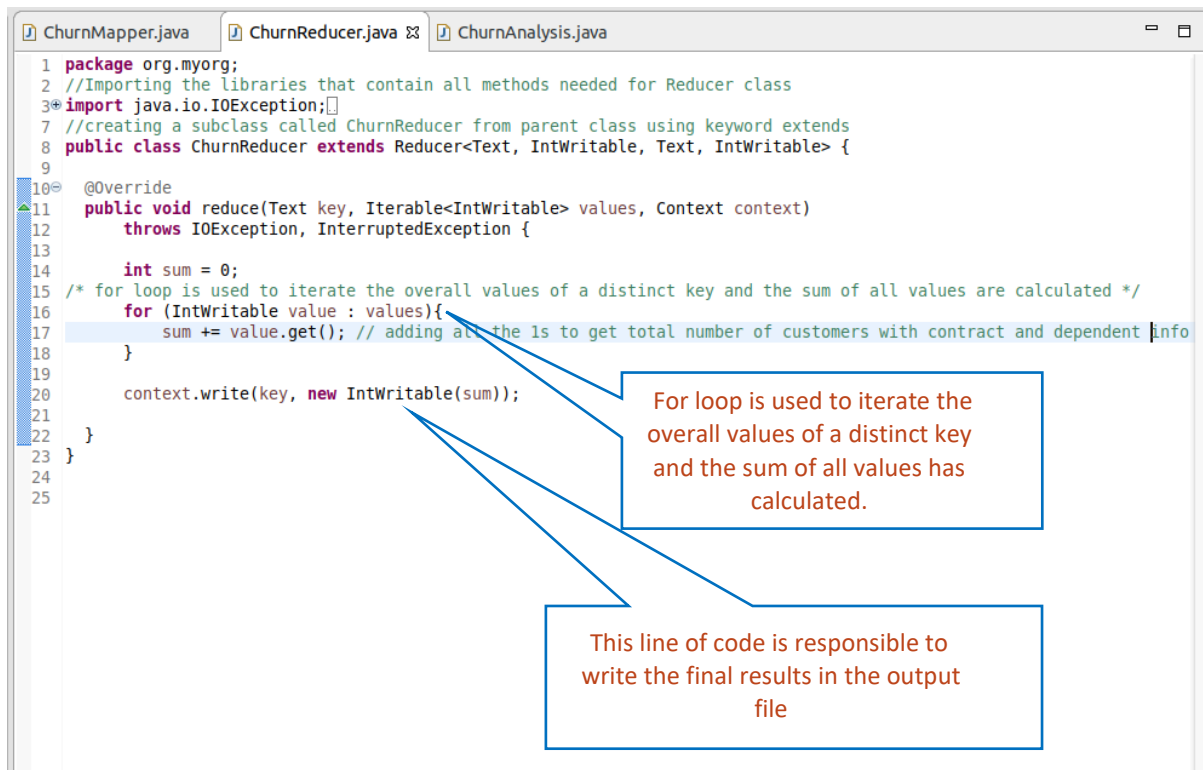


Figure 11

ChurnAnalysis (Driver Class):

This is the driver class of the program. It is implemented to run the whole project. Make sure both the input and output paths are passed in the command. It will create the job instance and is responsible for passing inputs to the mapper and reducer classes. It will delete the output folder from hdfs if that folder already exists. It returns a success message after the job gets completed successfully.

- **Job Name:** name of the job, job
- **Executable (Jar) Class:** the main executable class. Here it's, ChurnAnalysis.
- **Mapper Class:** class that overrides the "map" function. Here it's, ChurnMapper.
- **Reducer Class:** class that overrides the "reduce" function. Here it's, ChurnReducer.
- **Output Key:** type of output key. It's, Text.
- **Output Value:** type of output value. It's IntWritable.
- **File Input Path:** /Telecom/Input
- **File Output Path:** /Telecom/Output

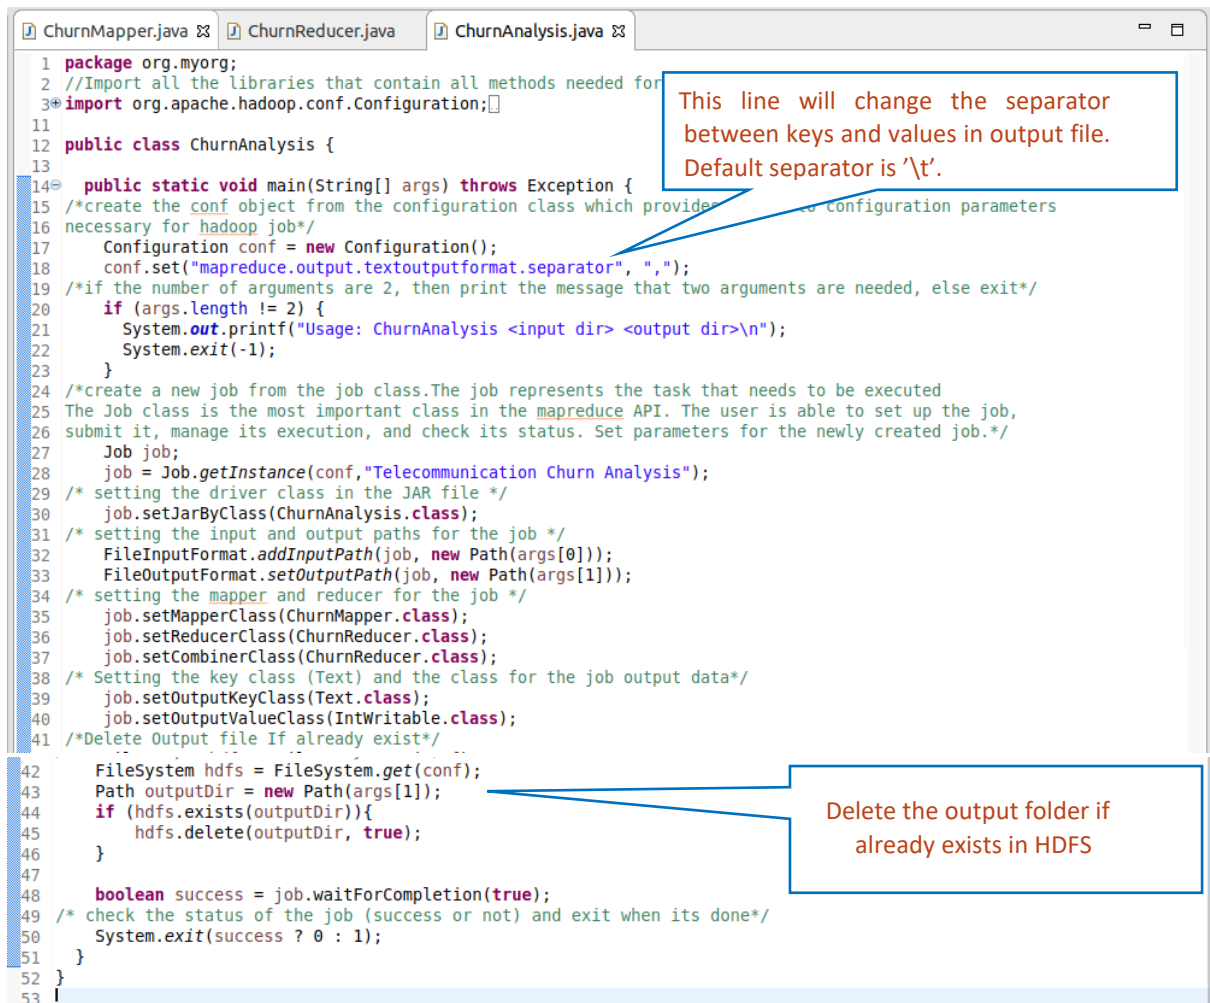
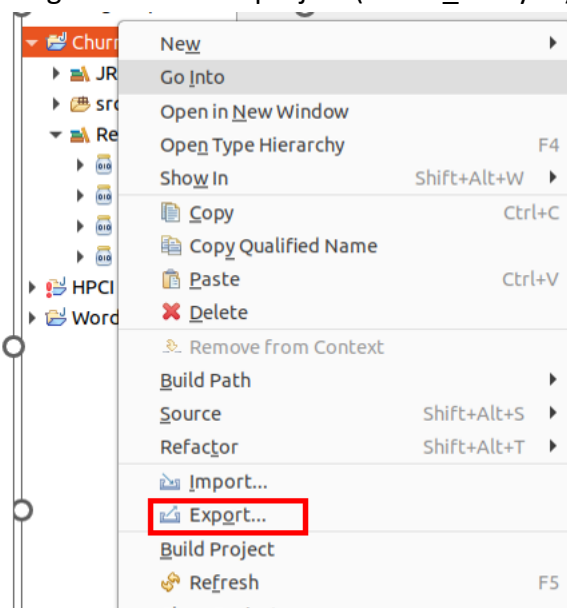



Figure 12

2.4 Export the Jar file

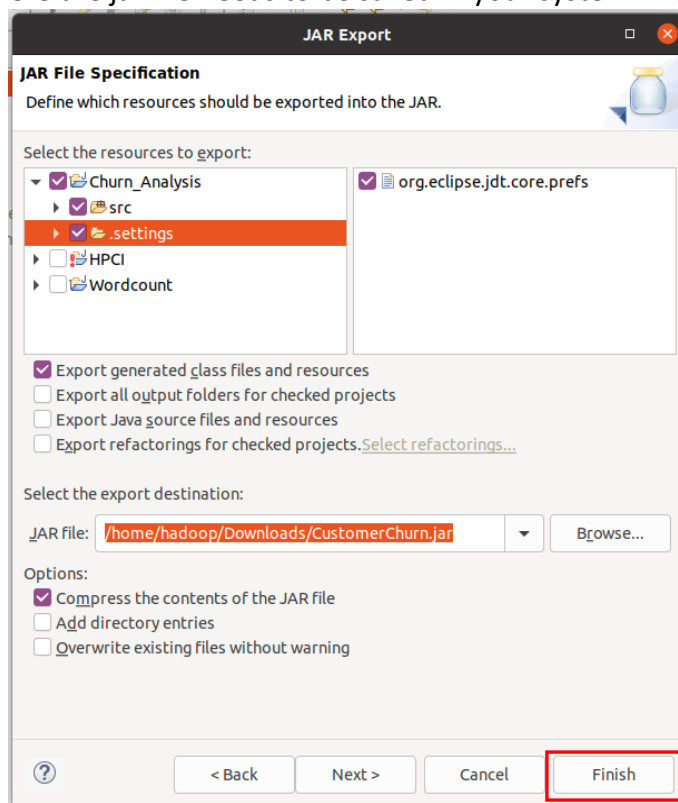
1. Right-click on the project (Churn_Analysis) and click on Export.



2. Click on the Next button



3. Select the resources to export (Churn_Analysis) and select the export destination where the jar file needs to be saved in your system. Hit the next button.



4. Select the main class and click on finish. It will export the jar in the given destination folder.

3. Results & Evaluation


1. The obtained output file 'part-r-00000' is in the format shown below.

```
(base) hadoop@hadoop-VirtualBox:~$ hdfs dfs -cat /Telecom/Output/part-r-00000
Month-to-month,No,No,1690
Month-to-month,No,Yes,1396
Month-to-month,Yes,No,530
Month-to-month,Yes,Yes,259
One year,No,No,825
One year,No,Yes,117
One year,Yes,No,482
One year,Yes,Yes,49
Two year,No,No,875
Two year,No,Yes,30
Two year,Yes,No,772
Two year,Yes,Yes,18
(base) hadoop@hadoop-VirtualBox:~$
```

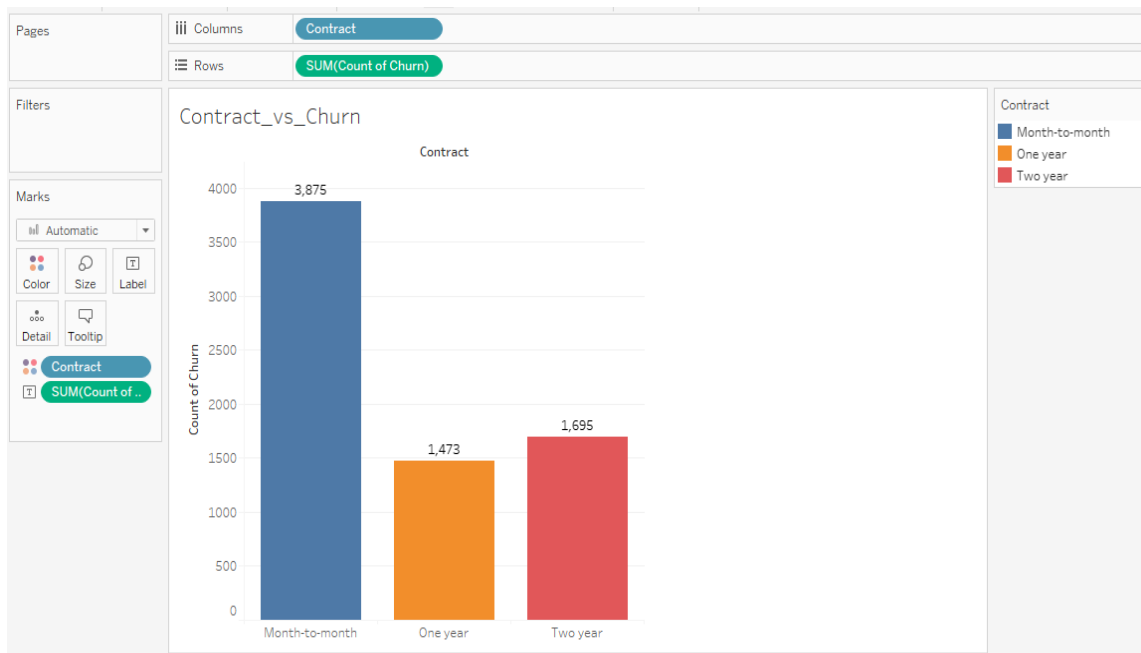
2. The output file has been converted to CSV and is imported into the tableau tool to visualize and identify the customer churned vary from contract period and people having dependents.

4. Observations:

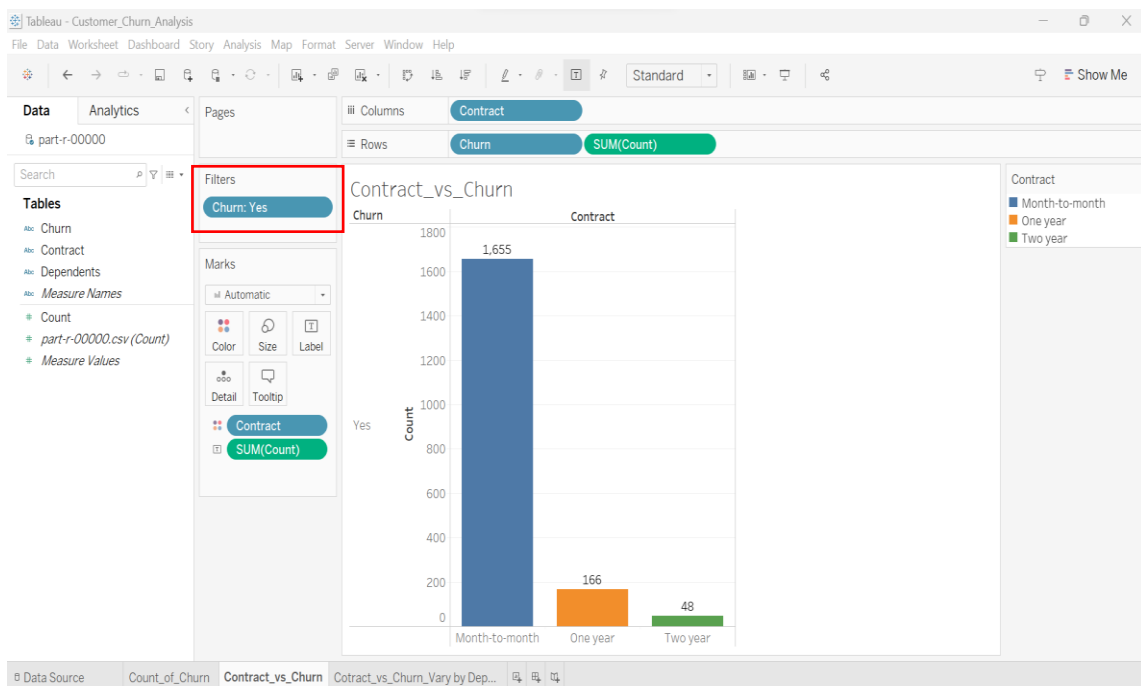
- a. It is observed that the company has lost around 25% of subscribers as shown in the plot below.



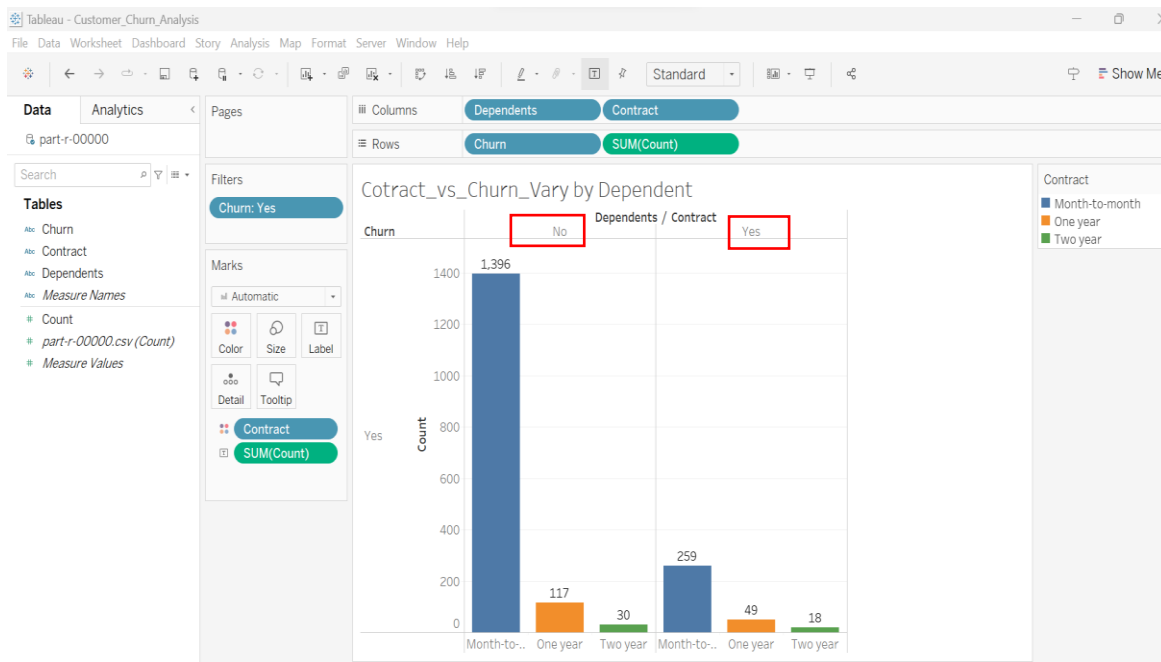
- b. It is observed that most of the subscribers like to choose month-to-month contracts as shown in the plot below.



c. It is observed that customers having month-to-month contract periods are more likely to churn compared to customers having one-year or two-year contract periods.



d. It is observed that the customers who do not have dependents and have subscribed for less contract period (month-to-month) are more likely to churn compared to the others as shown in the plot below.



Conclusion

With these observations, our research question is answered. It can be concluded that customers who subscribe for a month-to-month subscription, having no dependents, tend to cancel their subscription more frequently than those who opt for longer contract periods and have dependents. This difference is significant enough to affect the overall conclusion. Additionally, telecom companies can develop multiple strategies to improve customer retention and reduce churn. For example, companies can focus on offering long-term contract periods to customers, especially those without dependents. It not only develops their own business but also improves overall customer satisfaction.

Appendix

As part of the coursework, The Mapreduce_ChurnAnalysis Word document has been attached containing all supporting documents as an appendix to the coursework. It includes all the contents, such as the java classes code, the original dataset used for this project, the renamed dataset, the Log4j file, the output file, and a visualization created in Tableau (.twb file) to evaluate the results. This comprehensive list of contents in a single word document is intended to support and enhance the project's overall quality, allowing for a thorough analysis and evaluation of the project's findings.

References

hadoop.apache.org. (n.d.). Apache Hadoop 3.3.1 – Hadoop: Setting up a Single Node Cluster. [online] Available at: <https://hadoop.apache.org/docs/stable/hadoop-project-dist/hadoop-common/SingleCluster.html>.

hadoop.apache.org. (n.d.). MapReduce Tutorial. [online] Available at: https://hadoop.apache.org/docs/r1.2.1/mapred_tutorial.html.

Apache.org. (2015). Apache log4j 1.2 - Download Apache log4j 1.2. [online] Available at: <https://logging.apache.org/log4j/1.2/download.html> [Accessed 17 Mar. 2023].

help.tableau.com. (n.d.). Aggregate Functions in Tableau. [online] Available at:

https://help.tableau.com/current/pro/desktop/enus/calculations_calculatedfields_aggregate_create.htm [Accessed on 09 Mar, 2023].

logging.apache.org. (n.d.). LogManager (Apache Log4j API). [online] Available at: <https://logging.apache.org/log4j/2.x/log4j-api/apidocs/org/apache/logging/log4j/LogManager.html> [Accessed on 08 Mar, 2023].

Apache.org. (2015). Apache log4j 1.2 - Download Apache log4j 1.2. [online] Available at: <https://logging.apache.org/log4j/1.2/download.html>.

www.tutorialspoint.com. (n.d.). Eclipse - Create Jar Files. [online] Available at: https://www.tutorialspoint.com/eclipse/eclipse_create_jar_files.htm# [Accessed 17 Apr. 2023].

GeeksforGeeks. (2019). HDFS Commands. [online] Available at: <https://www.geeksforgeeks.org/hdfs-commands/> [Accessed on 08 Mar, 2023].